



NIPRGPT Frequently Asked Questions

Mr. Collen Roller

**Information Fusion Branch
Information Directorate (RI)
Air Force Research Laboratory**

**Version #: 1.4
Version Date: 11 September 2024**

Controlled by: AFRL/RIEA
Category(ies): CTI
Limited Dissemination Control: FEDCON
POC: Collen Roller

Purpose

Q: Why does the Department of Air Force need its own Large Language Model (LLM)?

A: The Department of the Air Force is **experimenting** with its relationship with Generative AI through the launch of NIPRGPT, an initiative in collaboration with the Air Force Research Laboratory. This experimental program enables Guardians, Airmen, civilian employees, and contractors to responsibly interact with Generative AI on the Non-classified Internet Protocol Router Network. NIPRGPT is unique compared to commercial tools as it is hosted on-premise in an Air Force-owned environment, allowing the DAF to use data on user interactions to perform research about Generative AI—something not possible with commercial tools. The initiative aims to investigate user experiences, interactions, and relationships with Generative AI, ensuring advanced technologies are provided at the speed of relevance to maximize competitive advantage. Adequate safeguards are in place to support this research as part of the DAF's ongoing modernization efforts.

Q: How long will NIPRGPT be available?

A: NIPRGPT is a backbone for GenAI research and will exist if it provides research value to advance and understand the role of GenAI to our mission. We expect as commercial tools are approved the use of NIPRGPT in those areas should diminish and other, new elements of the research would be advanced until the platform is no longer useful for that purpose.

Using NIPRGPT

Q: What type of information can I upload into the NIPRGPT RAG model?

A: You can upload any publicly available or Controlled Unclassified Data (CUI) data and International Traffic and Arms (ITAR). You should **not** upload Personal Identifiable Information (PII) or Protected Health Information (PHI).

Q: Can ITAR be uploaded?

A: Yes, you can upload ITAR.

Q: Can workspaces be public, or are they only private? In other words, can we make our workspace accessible to other NIPRGPT users?

A: Yes, workspaces can be both public and private. However, a public workspace, must be approved by the NIPRGPT team. You can create your own (private) workspace.

Q: Can someone upload a document to NIPRGPT directly?

A: Yes, you should use the RAG model.

Q: Do you have plans in the works for images, figures or graphs?

A: There are lots of features in the backlog, we have plans to build out the model in the future.

Q: Is there a limit to the amount of documents you can upload?

A: There is a 20 document limit.

Q: What is the size limit for uploads?

A: 40 Mb per upload

Getting Help and Providing Feedback

Q: How do I leave feedback?

A: Select the name icon in the upper right-hand corner of the page. Once the menu slides out, select the feedback icon which will display the feedback form.

Q: My documents won't upload, what do I do?

A: If you are struggling to upload documents, please reach out to our support email – afrl.ri.niprgpt@us.af.mil and provide the document that failed in addition to any error message received so we can investigate and identify the issue. If you're uploading multiple documents concurrently, we recommend limiting the documents you're uploading to 1-2 per upload.

Security and Accreditation

Q: What is the infrastructure that NIPRGPT resides in?

A: We are deployed using a hybrid infrastructure using both on-prem compute resources and cloud resources. The application sits on-prem and runs on our GPU cluster that we're currently building to support users for daily use and fine-tuning! Our cloud platform is called "Dark Saber". On-prem, we partner with the HPC (part of the DoD HPC) and are currently working hard to scale our system.

Q: Does NIPRGPT use Platform One?

A: We partner with Platform One and leverage their code-as-infrastructure, but we do not run on Platform One. Dark Saber is a separate platform. Initially, we were using Platform One's DNS but we've migrated away from that to <https://niprgpt.mil> over the past few months.

Q: Does it connect to ChatGPT?

A: No! Part of our approved boundary is that all inference is done within our boundary (on-prem and or GovCloud). Models are hosted on-prem and do not reach out to other locations.

Q: Is it secluded and locked for government use only or does it touch the internet?

A: Users can access NIPRGPT from the commercial internet but can only access the platform with a verified .mil email address. We have a ZeroTrust access control system mirrored from PlatformOne. In addition, we support login using the DoD Global Directory for single-sign-on.

Q: What is the accreditation status?

A: It has an Authority to Operate (ATO) at the IL5 level from the Air Force Research Laboratory (AFRL).

Data Management

Q: How is the data tagged or separated?

A: Data is separated on an individual basis. Each user has workspaces they can create and upload documents to interact with. Each of these workspaces are accessible only by the user who created the workspace. We do have "shared" workspaces but these are controlled by the NIPRGPT team to prevent potential leaks.

Technical Details

Q: Can NIPRGPT be pointed at a specific data lake that we want to interact with?

A: Not yet, stay tuned, APIs are coming.

Q: Will NIPRGPT work on DREN and NIPR?

A: Yes, it should work on both.

Q: How do I get access to the code?

A: NIPRGPT is a GOTS solution that can support the entire DoD. If you're interested in using the code to deploy your own Generative AI solution, please reach out to the team, afrl.ri.niprgpt@us.af.mil, and we will provide you access to our repository.

Q: How do I get an air-gapped deployment of NIPRGPT?

A: NIPRGPT is a GOTS solution that you can use at your own respective shop. However, it will require access to GPUs to optimally run the Large Language Model. Most locations that want a deployment are trying to support ~4000 users. For this, we recommend 4 NVIDIA H100 GPUs connected together using an NVIDIA NVL or 8 NVIDIA L40S GPUs. For more specifics, reach out to us and we can help provide a better estimate based on your specific requirements.

Contributing and Supporting

Q: I want to help, how do I contribute?

A: The team needs the following:

1. Software developers who have foundations in Large Language Models, Python, RAG and want to support Generative AI for the DoD.
2. Infrastructure support for AWS, Google, Terraform, Kubernetes and Docker.
3. Web development support with experience in ReactJS and ExpressJS APIs.

Q: How can I best support my units use of NIPRGPT?

A: Reach out to us at afrl.ri.niprgpt@us.af.mil.

Troubleshooting

Q: Why am I getting a “504 Bad Gateway” message?

A: Try refreshing the page or opening a new browser tab.

Q: Why am I getting a “Failed to create a workspace” or why is my workspace not loading?

A: Try logging out and then logging back in.